



A FRAMEWORK TO SUPPORT TEACHERS IN IDENTIFYING KNOWLEDGE GAPS

EPFL PhD Summit

Advisor: Carlos Eduardo Pedreira

Laura de Oliveira F. Moraes

About me



- PhD Student at UFRJ, Brazil
 - *Research field:* Artificial Intelligence, Education.
 - Machine Teaching: This project aims to use elements on how machines learn to improve how students learn.
- Software Engineer at CERN for 4 years
- 2018 Fellow at Data Science for Social Good
- Data Science startup co-founder (TWIST Systems)

Motivation

- Understand gaps in knowledge **continuously**
 - Students are evaluated only in specific times (pre-exam anxiety)
 - Improve teaching methods
 - Personalized education
- Not domain specific
- Python language
 - Fastest-growing major programming language



Goal

Assess students' knowledge continuously, avoiding pre-exam anxiety and late diagnosis.

- ★ Input: questions, concepts and a student
- ★ Output: student knowledge inference

Welcome to the Machine Teaching Project!

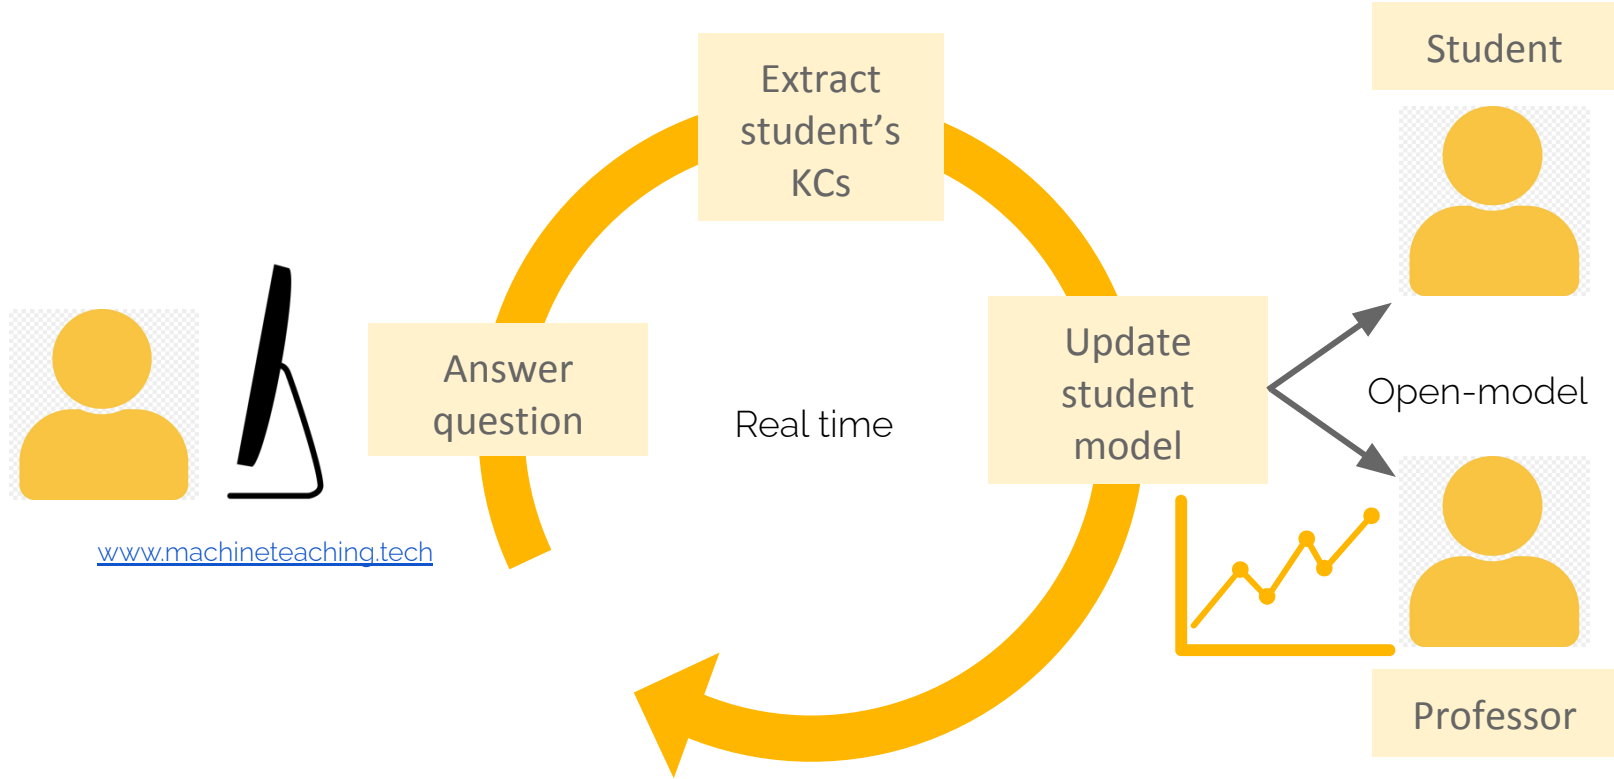
Start 



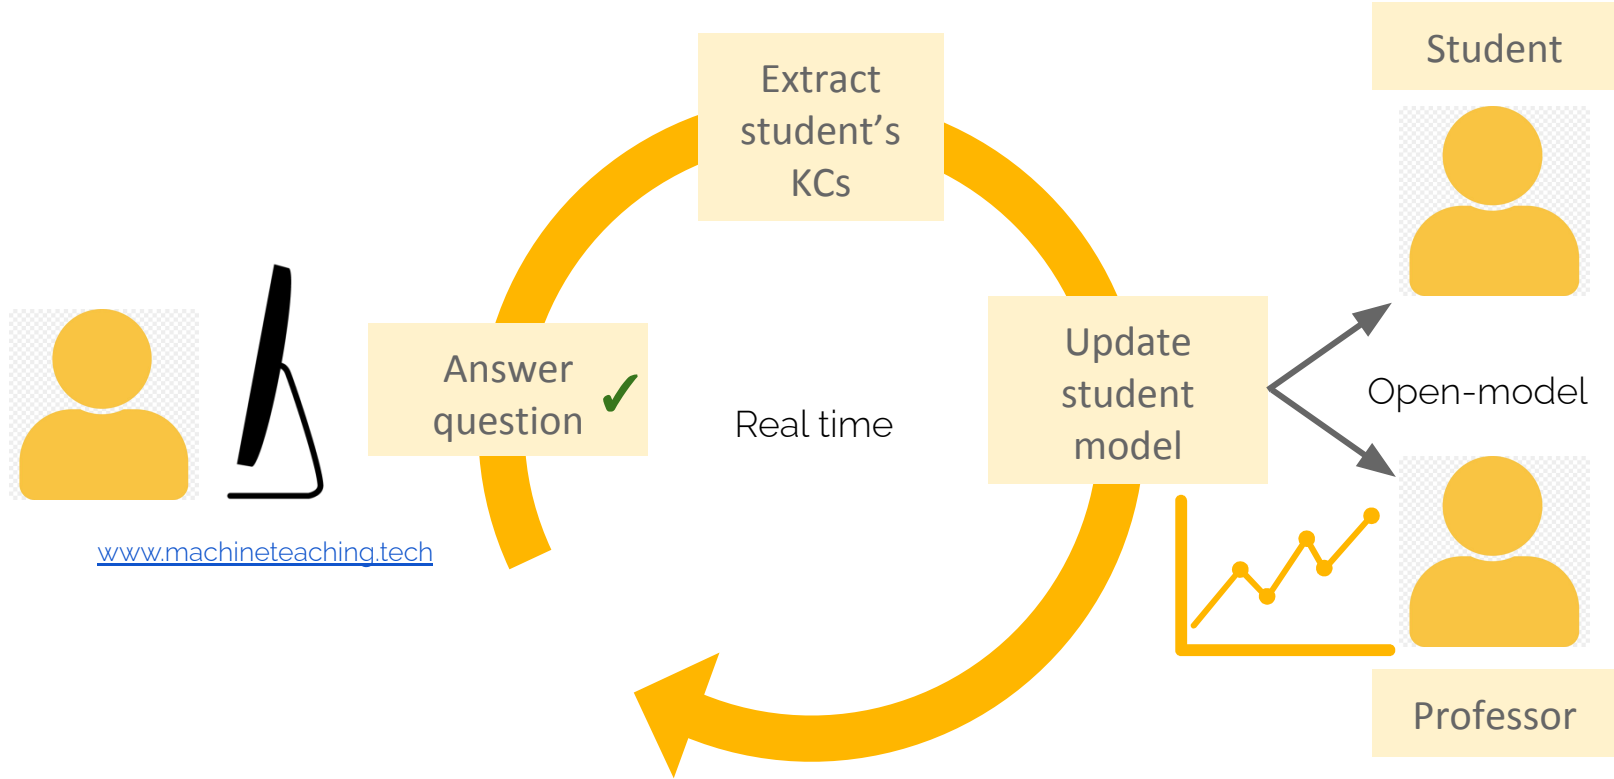
Contributions so far...

- ✓ An open-source anonymized database:
 - a. **435 college freshman** with at least one submission
 - b. **854** different **problems** (90 revised)
 - c. **21,884** different **submissions** (in 5,607 problems)
 - d. Running this semester in **4 classes**
- ✓ Unsupervised knowledge component (KC) discovery
- ✓ Personalized student model.
- ✓ Real user studies.
- ✓ Integrated framework

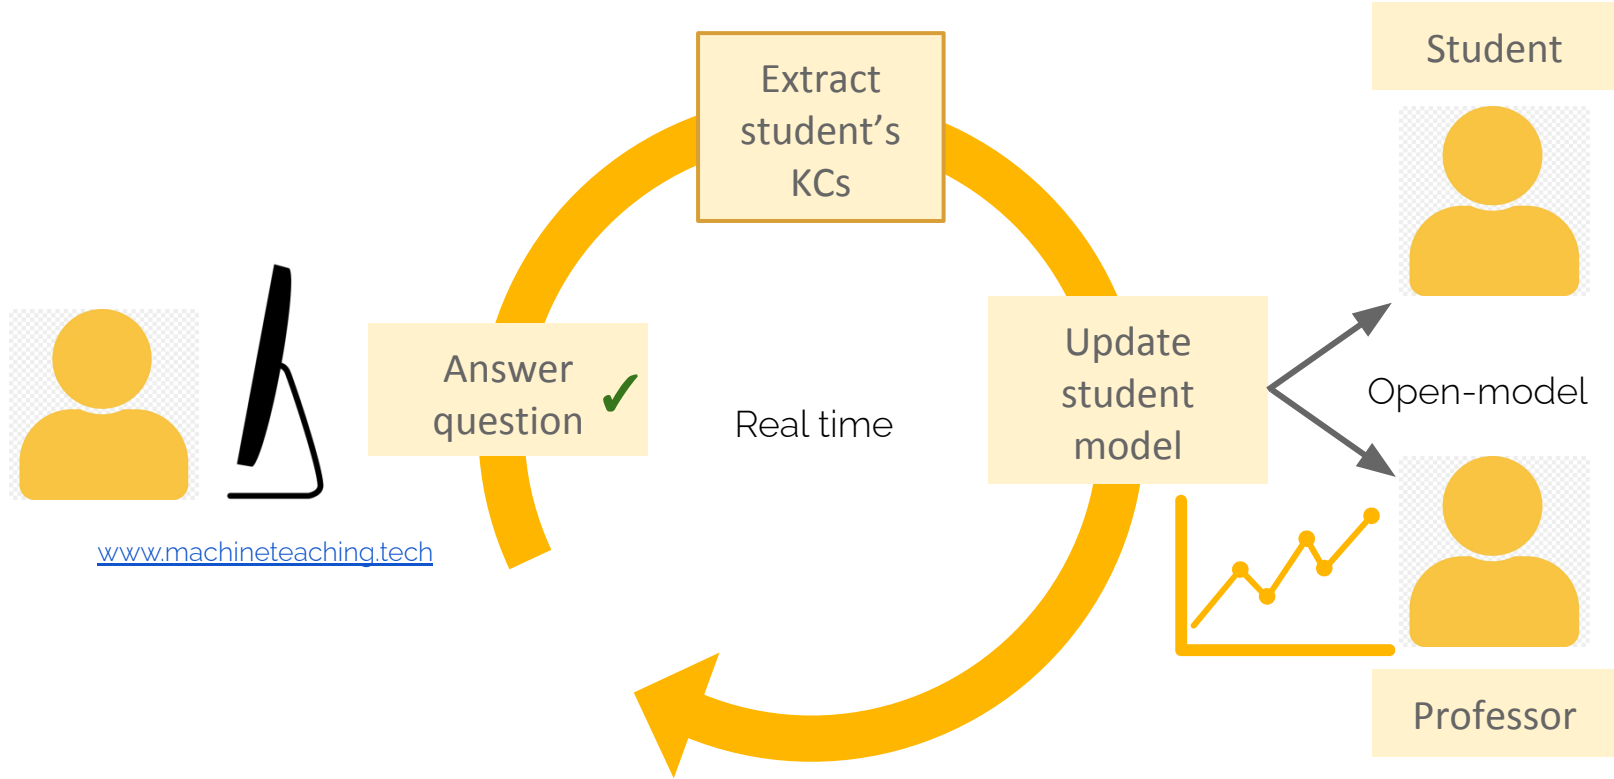
How does it work?





How does it work?





How does it work?



Methodology



Methodology



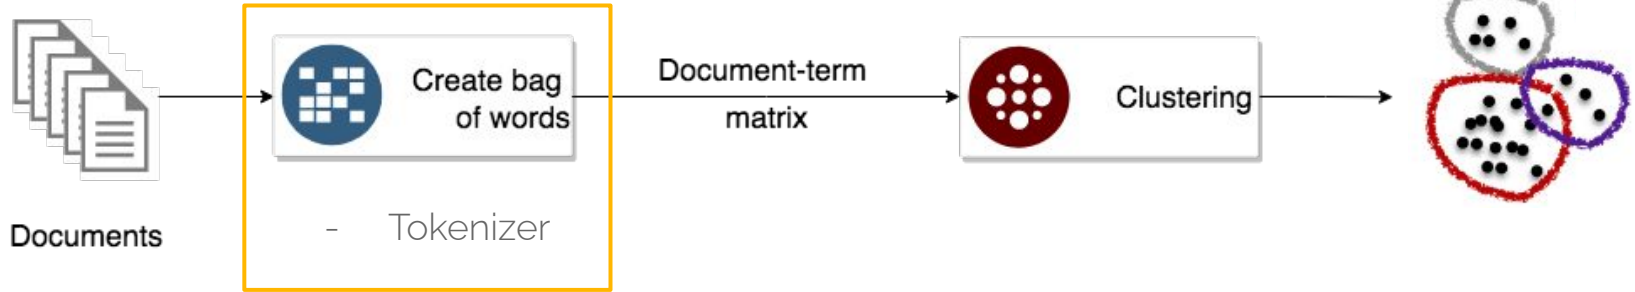

Document-term matrix

documents
code snippets
answers

| | for | if | append |
|---|-----|----|--------|
| 1 | 2 | 1 | 1 |
| 2 | 0 | 1 | 1 |
| 3 | 1 | 2 | 0 |

terms/tokens/words

Methodology



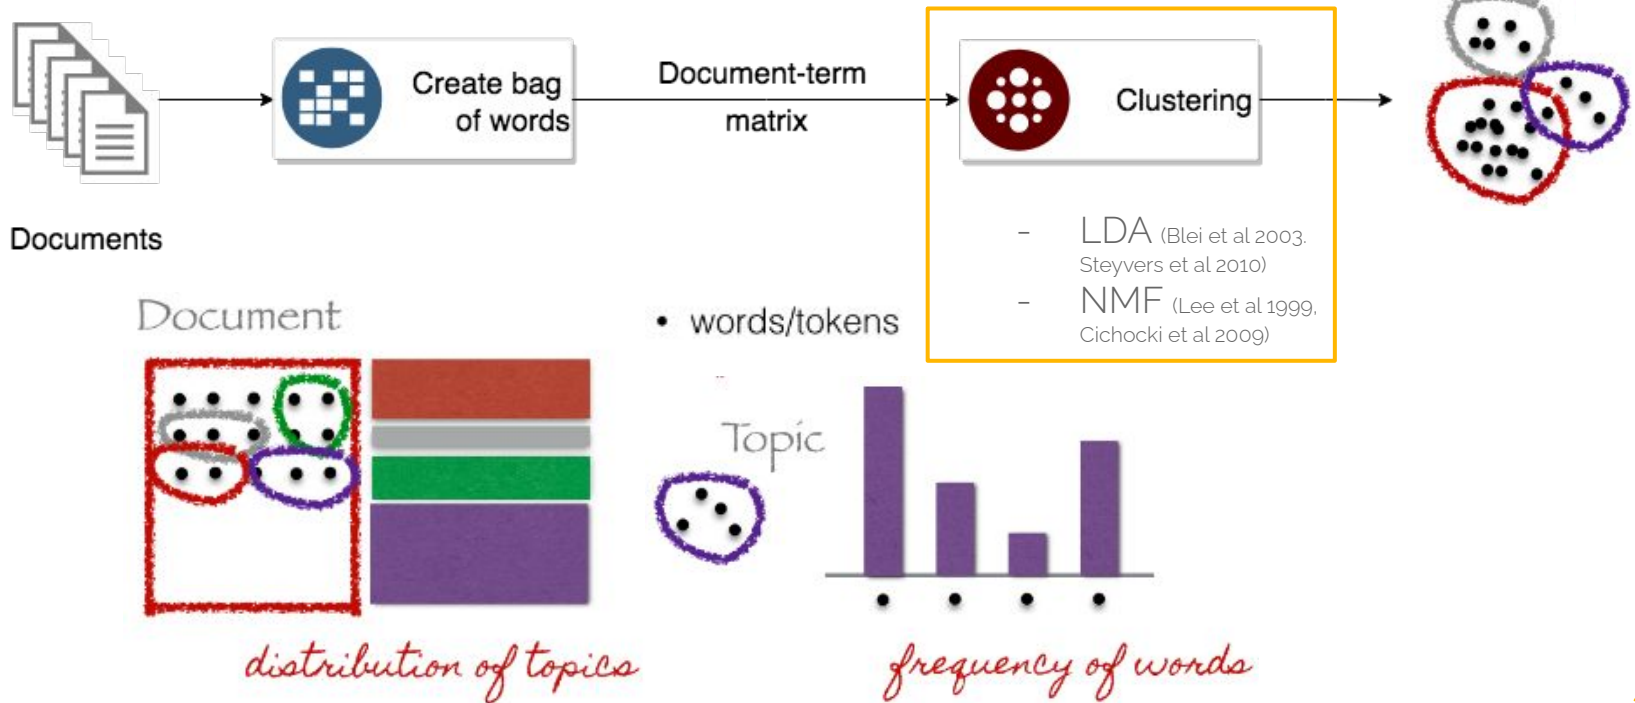

Document-term matrix

documents
code snippets
answers

| | for | if | append |
|---|-----|----|--------|
| 1 | 2 | 1 | 1 |
| 2 | 0 | 1 | 1 |
| 3 | 1 | 2 | 0 |

terms/tokens/words

Methodology





Given all these possibilities, how to choose the best set of hyper-parameters, the clustering method and the number of clusters?



Evaluation - Topic Coherence

(Mimno et al 2011, Röder et al 2015)

- **Human interpretability** metric
- Ratio between **co-occurrence of top-N** words and their **total occurrence** within the topic.
- Same words appear in the same documents of the topic.

$$C_{UMass}(t_k, V) = \sum_{m=1}^N \sum_{j=1}^N \log \frac{Q(v_j, v_m, t_k) + \epsilon}{Q(v_j, t_k)}$$

Methodology



Grid search (Bergstra et al 2012)

$$\begin{array}{c} 10 \\ \text{Minimum Document Frequency Values} \end{array} \times \begin{array}{c} 2 \\ \text{Binary Appearances Options} \end{array} \times \begin{array}{c} 3 \\ \text{Vectorizers} \end{array} \times \begin{array}{c} 2 \\ \text{Clustering Methods} \end{array} \times \begin{array}{c} 14 \\ \text{Number of Clusters/Topics} \end{array} = 1680 \text{ RUNS}$$

Dataset



```
def square(number):  
    square_dict=dict()  
    for i in range(1,number+1):  
        square_dict[i]=i*i  
    return square_dict
```

```
def factorial(number):  
    total = 1  
    for i in range(number, 1, -1):  
        total = total * i  
    return total
```


54 revised Python snippets
to be used as **documents**
(now, there are 90)

Taking a deeper look

| Experiment Id | Min DF | Binary | Vectorizer | Method | Best-k |
|---------------|-------------|-------------|--------------|------------|-----------|
| 26 | 0.05 | True | Count | LDA | 12 |

Document distribution per topic:

- Topic 1: 2
- Topic 2: 1
- Topic 3: 1
- Topic 4: 2
- Topic 5: 1
- **Topic 6: 13**
- Topic 7: 1
- **Topic 8: 14**
- Topic 9: 0
- **Topic 10: 7**
- Topic 11: 0
- **Topic 12: 12**



Turning Topics into Concepts



String
manipulation

7. Data type: string

14. Function

8. Data type: array or list

Conditional
structure

6. Logic 11. Conditional

14. Function

Math and
string loops

5. Math

7. Data type: string

12. Loop

Math
functions

1. Syntax

3. Data type: number

5. Math 14. Function

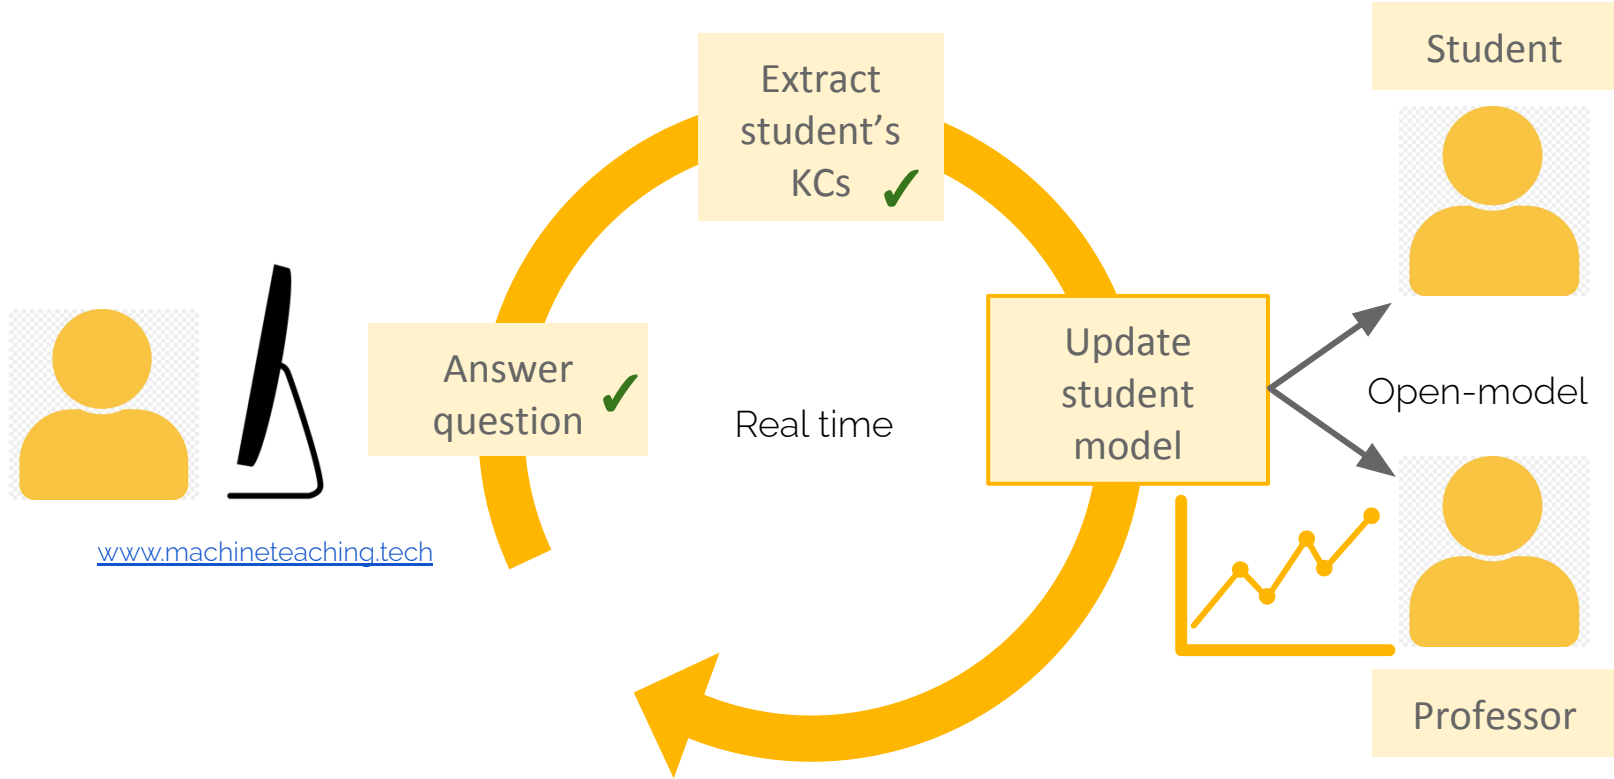
List loops

8. Data type: list or array

11. Conditional 12. Loop

14 professors
evaluated the
results in 3
tasks

Review



Student models

- Bayesian Knowledge Tracing (Corbett et. al. 1995)
 - Been widely used
 - Several enhancement proposals
- Performance Factor Analysis (Pavlik et al 2009)
- Tensor factorization (Sahebi et al 2016)




Figure reproduced from Sahebi et al 2016

Student Performance Kit

- Python library like Scikit-Learn
 - 3 models: **BKT**, PFA, TF
 - Scores: **LL**, **AIC**, **BIC**, **RMSE**, **Acc**, AUC

spkit 0.0.1

```
pip install -i https://test.pypi.org/simple/ spkit==0.0.1
```



Conclusions

- **Engage** a team of students, professors and teaching assistants.
- Novel approach to **code clustering** in the EDM context.
- Good clustering schemes suited for **human interpretability**.
- The majority of the methodology can be **generalized to other domains**.
- Open and reproducible (still needs a bit of organization!):
 - <https://github.com/laura-moraes/machine-teaching>
 - <https://github.com/lauramoraes/StudentPerformanceKit>

Next Questions

- How much domain knowledge and specialization is needed to transpose it to another domain, such as History or Social Studies?
- Can we recommend personalized content in order to improve learning?



Thanks!

Any questions?



Students complain a little about the site, but it is because they are not very aware of how useful it is to themselves. For me the level of the practical classes has greatly improved.

Hugo Nobrega, professor

3

Additional material

Related Work

- Autograders
-
- **Question classification**
 - Specialists
 - Supervised Learning
 - Unsupervised Learning
 - ASTs
 - **Matrix factorization**

CS1 Concepts



4 References were used to define the concepts usually given in a CS1 course:

- ACM Computer Science Curricula 2013
- Sheard et al 2011
- Petersen et al 2011
- Cherenkova et al 2014

LIST OF FINAL CONCEPTS:

- | | | | |
|-----|--------------------------|-----|-------------|
| 1. | Syntax | 11. | Conditional |
| 2. | Assignment | 12. | Loop |
| 3. | Data type: number | 13. | Nested loop |
| 4. | Data type: boolean | 14. | Function |
| 5. | Math | 15. | Recursion |
| 6. | Logic | | |
| 7. | Data type: string | | |
| 8. | Data type: list or array | | |
| 9. | Data type: tuple | | |
| 10. | Data type: dict | | |



Create bag
of words



Tokenizer

Separate text in:

- **Keywords**
- **Operators**
- **Data types**
- **Indents/Dedents**

Minimum
Document
Frequency

Choose a value in a
range from 0.05 to 0.5
(percentage of
documents)

Binary
Appearance


How to count word
appearance:

1. Just once per document **OR**
2. Each time it appears in the document

Vectorizer

How to count frequency
of words:

1. Regular Count **OR**
2. Tf-Idf weights
(Salton et al 1986, Zhang et al 2011)
OR
3. NCut weights
(Yan et al 2012)



Latent Dirichlet Allocation (**LDA**)

(Blei et al 2003, Steyvers et al 2010)

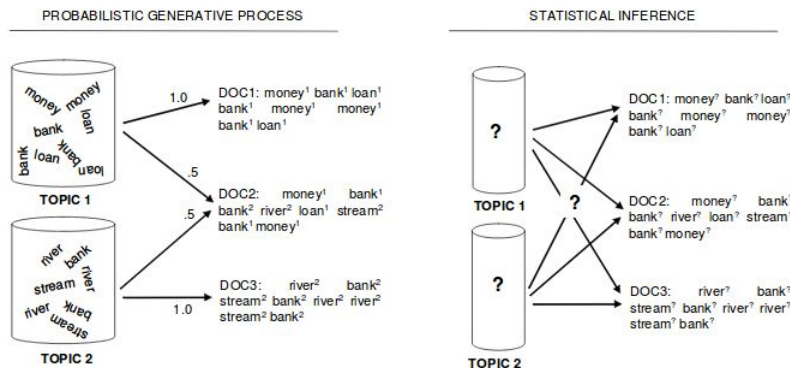


Figure reproduced from Steyvers et al 2010

Non-negative Matrix Factorization (**NMF**)

(Lee et al 1999, Cichocki et al 2009)




Figure reproduced from <http://tech.opentable.com/2015/01/12/finding-key-themes-from-free-text-reviews/>



Experiments

Objective: group code snippets used to solve CS1 exercises according to the concepts used to solve the problem.

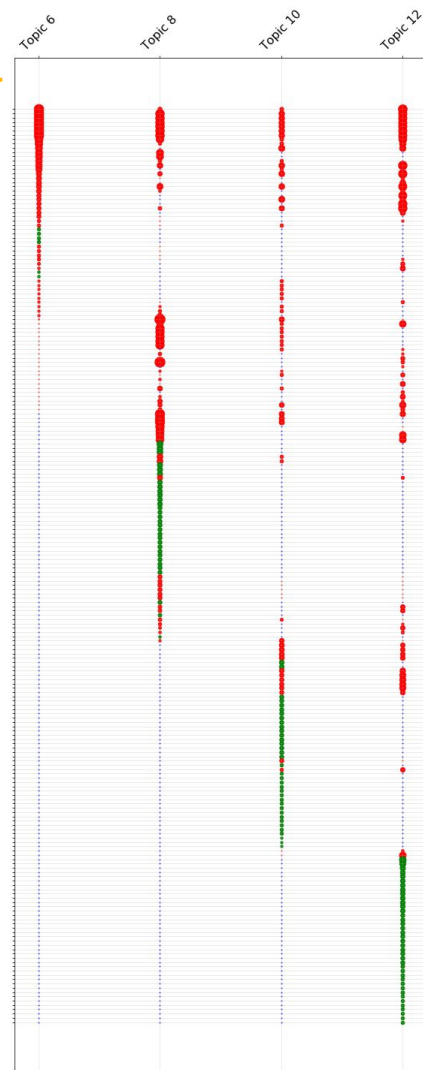
Hypothesis: the words used to code a CS1 exercise are an indicative of the concepts needed to solve them.

Methodology: cluster these code snippets using as bag of words the words used in the code and the presented methodology.

Taking a deeper look




The main topics 6, 8, 10 and 12 correspond to 85% of the documents and 77.4% of the terms.




Topic 8

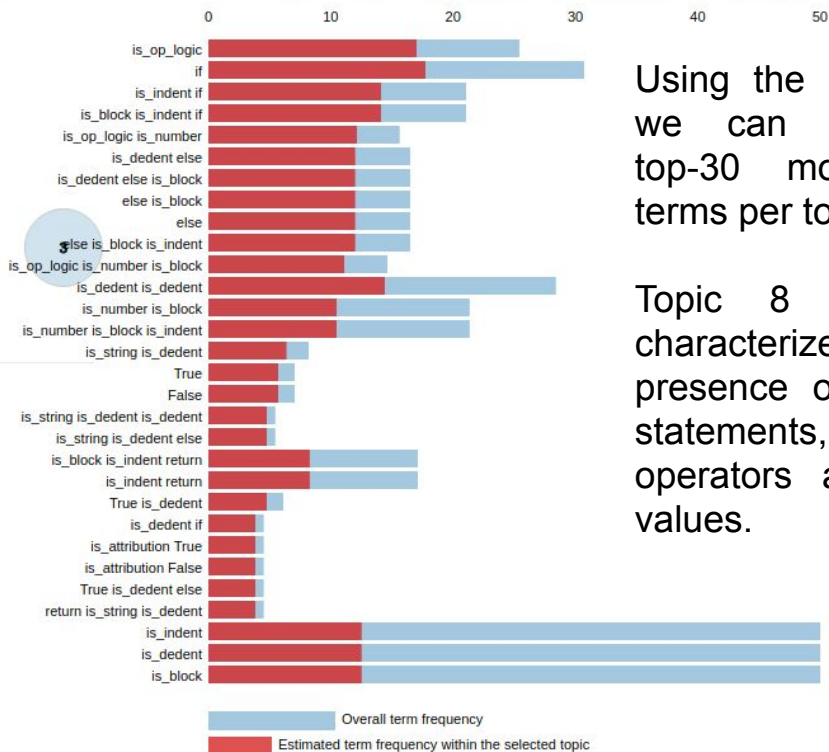
Intertopic Distance Map (via multidimensional scaling)



Marginal topic distribution



Top-30 Most Relevant Terms for Topic 8 (25.9% of tokens)



Using the LDAVis tool, we can inspect the top-30 most relevant terms per topic.

Topic 8 is strongly characterized by the presence of conditional statements, logical operators and boolean values.

1. $\text{saliency}(\text{term } w) = \text{frequency}(w) * [\sum_t p(t|w) * \log(p(t|w)/p(t))]$ for topics t ; see Chuang et. al (2012)
 2. $\text{relevance}(\text{term } w | \text{topic } t) = \lambda * p(w|t) + (1 - \lambda) * p(w|t)/p(w)$; see Sievert & Shirley (2014)

Topic 8: code snippets

```
def is_prime(number):  
    '''Returns True for prime numbers, False otherwise'''  
    #Edge Cases  
    if number == 1:  
        prime = False  
    elif number == 2:  
        prime = True  
    #All other primes  
    else:  
        prime = True  
        for check_number in range(2, int(number/2)+1):  
            if number % check_number == 0:  
                prime = False  
                break  
    return prime
```


```
def light(switchA, switchB):  
    if switchA == 1 and switchB == 1:  
        return True  
    else:  
        return False
```

```
def max_of_three(a,b,c):  
    max_number = 0  
    if a > b:  
        if a > c:  
            max_number = a  
        else:  
            max_number = c  
    else:  
        if b > c:  
            max_number = b  
        else:  
            max_number = c  
    return max_number
```

```
def days_in_month(month, year):  
    if month == 2:  
        if (year % 400) == 0:  
            return 29  
        elif (year % 100) == 0:  
            return 28  
        elif (year % 4) == 0:  
            return 29  
        else:  
            return 28  
    elif month in (4,6,9,11):  
        return 30  
    elif month in (1,3,5,7,8,10, 12):  
        return 31
```

Topic 12

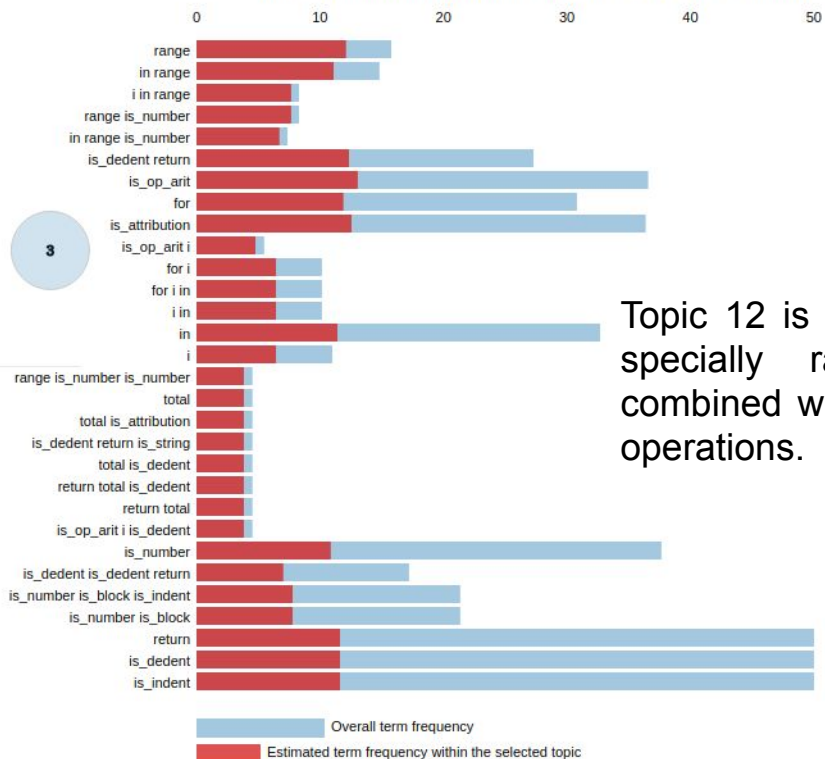
Intertopic Distance Map (via multidimensional scaling)



Marginal topic distribution



Top-30 Most Relevant Terms for Topic 12 (24.6% of tokens)



Topic 12 is about loops, specially range loops combined with arithmetic operations.

1. saliency(term w) = frequency(w) * [sum_t p(t | w) * log(p(t | w)/p(t))] for topics t; see Chuang et. al (2012)

2. relevance(term w | topic t) = λ * p(w | t) + (1 - λ) * p(w | t)/p(w); see Sievert & Shirley (2014)

Topic 12: code snippets

```
def digit_sum(digit):  
    total = 0  
    for i in range(1,5):  
        number = "%5" % digit  
        number = int(number * i)  
        total = total + number  
    return total
```


```
def fatorial(number):  
    total = 1  
    for i in range(number, 1, -1):  
        total = total * i  
    return total
```

```
def all_even():  
    values = []  
    for a in range(0,9,2):  
        for b in range(0,9,2):  
            for c in range(0,9,2):  
                values.append('2'+str(a)+str(b)+str(c))  
    return ", ".join(values)
```


```
def binby5(number_csv):  
    value = []  
    items=[x for x in number_csv.split(',')]  
    for p in items:  
        intp = int(p, 2)  
        intp = int(p[0])*8 + int(p[1])*4 + int(p[2])*2 + int(p[3])  
        if not intp%5:  
            value.append(p)  
    return ", ".join(value)
```

Topic 10

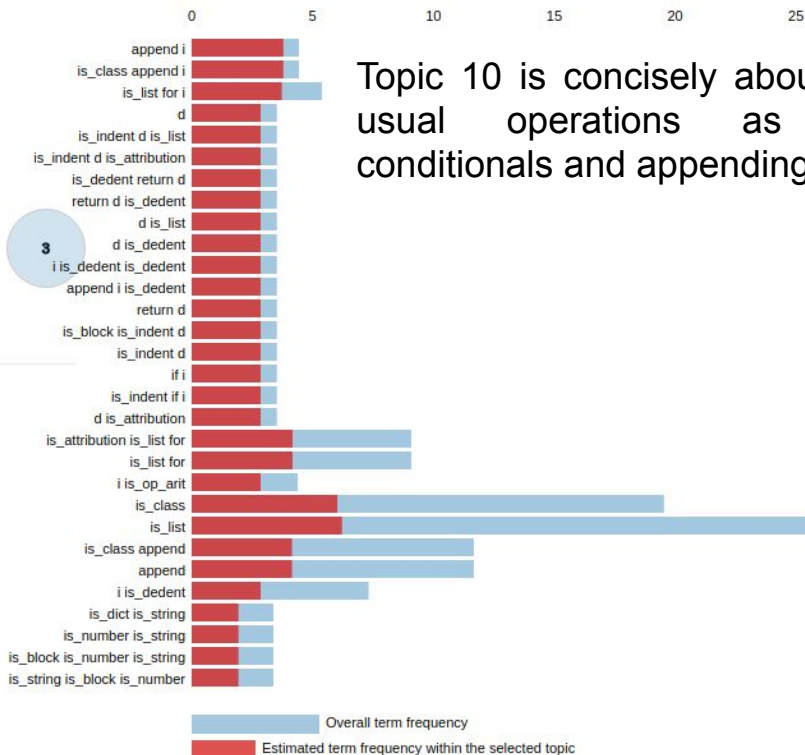
Intertopic Distance Map (via multidimensional scaling)



Marginal topic distribution



Top-30 Most Relevant Terms for Topic 10 (14.1% of tokens)



Topic 10 is concisely about lists and its usual operations as for loops, conditionals and appending elements.

1. $\text{saliency}(\text{term } w) = \text{frequency}(w) * [\sum_t p(t | w) * \log(p(t | w) / p(t))]$ for topics t ; see Chuang et. al (2012)
 2. $\text{relevance}(\text{term } w | \text{topic } t) = \lambda * p(w | t) + (1 - \lambda) * p(w | t) / p(w)$; see Sievert & Shirley (2014)

Topic 10: code snippets

```
def common(list1, list2):
    common_list = []
    for i in list1:
        if i in list2 and i not in common_list:
            common_list.append(i)
    common_list.sort()
    return common_list
```


```
def count(sentence):
    d={"digits":0, "letters":0}
    for char in sentence:
        if char.isdigit():
            d["digits"]+=1
        elif char.isalpha():
            d["letters"]+=1
        else:
            pass
    return d
```

```
def dedupe(dup_list):
    nodup_list = []
    for i in dup_list:
        if i not in nodup_list:
            nodup_list.append(i)
    return nodup_list
```


```
def divisible():
    l=[]
    for i in range(2000, 3201):
        if (i%7==0) and (i%5!=0):
            l.append(i)
    return l
```


Topic 6

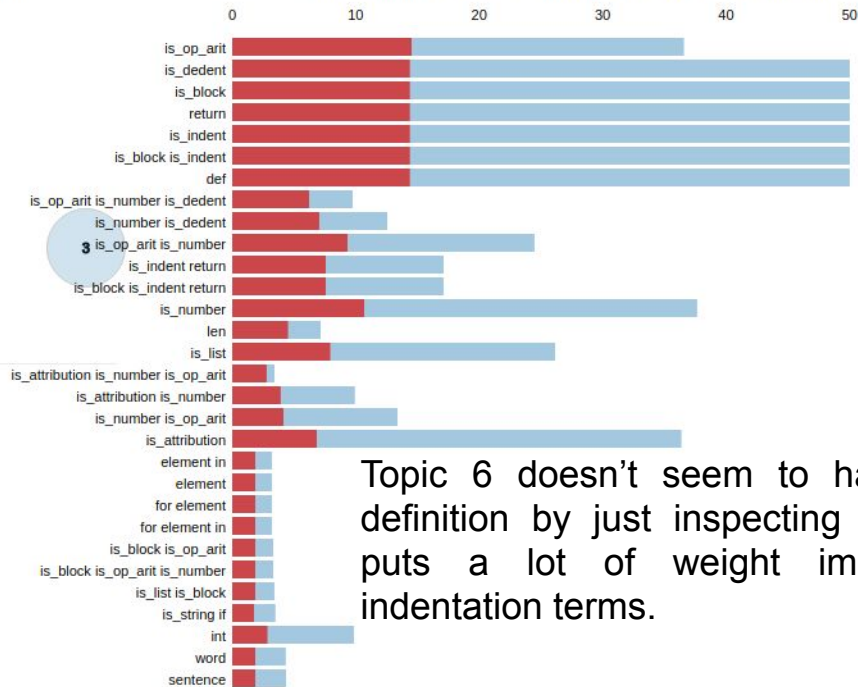
Intertopic Distance Map (via multidimensional scaling)



Marginal topic distribution



Top-30 Most Relevant Terms for Topic 6 (12.8% of tokens)



Topic 6 doesn't seem to have a clear definition by just inspecting its terms. It puts a lot of weight importance in indentation terms.

Overall term frequency
Estimated term frequency within the selected topic

1. saliency(term w) = frequency(w) * $[\sum_t p(t | w) * \log(p(t | w) / p(t))]$ for topics t ; see Chuang et. al (2012)

2. relevance(term w | topic t) = $\lambda * p(w | t) + (1 - \lambda) * p(w | t) / p(w)$; see Sievert & Shirley (2014)

Topic 6: code snippets

```
import math
def formula(D):
    C = 50
    H = 30
    Q = round(math.sqrt(2*C*D/float(H)))
    return Q
```

```
def sum_str(s1,s2):
    return int(s1)+int(s2)
```

```
def square(num):
    return num ** 2
```

By analyzing the code snippets, topic 6 comprises codes with one indentation structure (simple coding structures), sometimes without even the need to assign variables to solve the exercise.

```
def euro_conversion(amount, exchange_rate):
    euro = int(amount//exchange_rate)

    euro50s = int(euro // 50)
    remainingEuros = euro % 50

    euro20s = int(remainingEuros // 20)
    remainingEuros = remainingEuros % 20

    euro10s = int(remainingEuros // 10)
    remainingEuros = remainingEuros % 10


    euro5s = int(remainingEuros // 5)
    remainingEuros = remainingEuros % 5

    return(euro, euro50s, euro20s, euro10s, euro5s, remainingEuros)
```

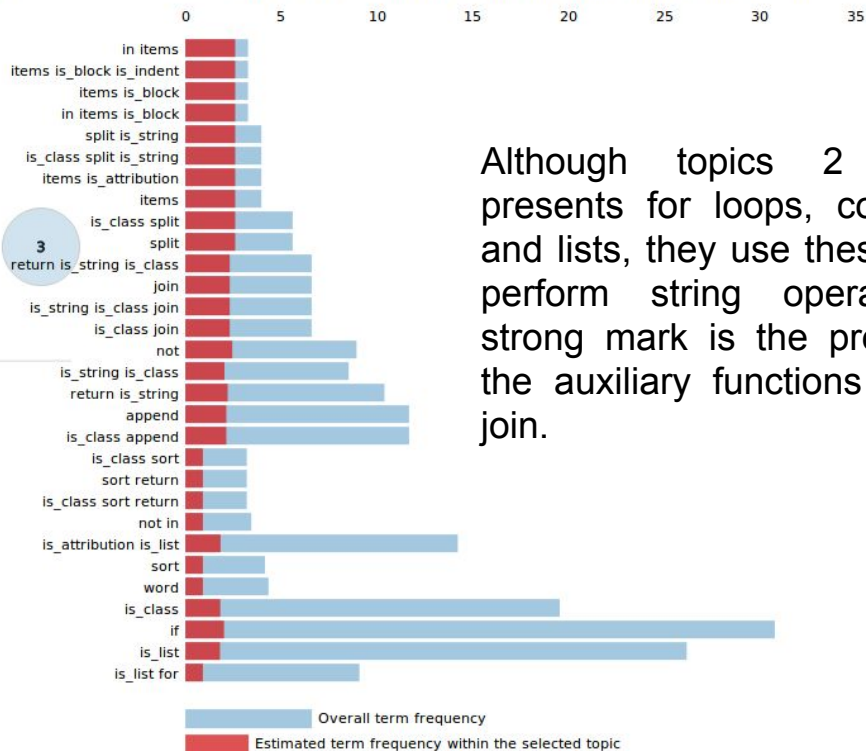
Topics 2 and 4

Intertopic Distance Map (via multidimensional scaling)

Top-30 Most Relevant Terms for Topic 4 (4.7% of tokens)



Marginal topic distribution



Although topics 2 and 4, presents for loops, conditionals and lists, they use these tools to perform string operations. A strong mark is the presence of the auxiliary functions split and join.

1. $saliency(\text{term } w) = \text{frequency}(w) * [\sum_t p(t | w) * \log(p(t | w) / p(t))]$ for topics t ; see Chuang
 2. $relevance(\text{term } w | \text{topic } t) = \lambda * p(w | t) + (1 - \lambda) * p(w | t) / p(w)$; see Sievert & Shirley (2014)

Topics 2 and 4: code snippets

```
def check_password(passwords):  
    import re  
    value = []  
    items=[x for x in passwords.split(',')]  
    for p in items:  
        if len(p)<6 or len(p)>12:  
            continue  
        else:  
            pass  
            if not re.search("[a-z]",p):  
                continue  
            elif not re.search("[0-9]",p):  
                continue  
            elif not re.search("[A-Z]",p):  
                continue  
            elif not re.search("[$#@]",p):  
                continue  
            elif re.search("\s",p):  
                continue  
            else:  
                pass  
            value.append(p)  
    return ",".join(value)
```

```
def sort_dedupe(words):  
    items = words.split(' ')  
    items_dedupe = []  
    for word in items:  
        if word not in items_dedupe:  
            items_dedupe.append(word)  
    items_dedupe.sort()  
    return " ".join(items_dedupe)
```





```
def sort_csv(csv):  
    items = csv.split(',')  
    items.sort()  
    return ",".join(items)
```

Turning Topics into Concepts




- **Theme identification:** the professors were shown four codes from the same topic. They should **label each topic** with a simple description.
- **Concept identification:** each professor should **associate up to three concepts** (from the 15 available) to each presented code.
- **Intruder identification:** the professors should **identify the intruder document given a topic.**

Theme identification



- Topic 4: String manipulation
- Topic 6: Math functions
- Topic 8: Conditional structure
- Topic 10: List loops
- Topic 12: Math and string loops

Intruder identification



References

ACM Computer Science Curricula 2013

Bergstra, J., & Yoshua Bengio, U. (2012). Random Search for HyperParameter Optimization. Journal of Machine Learning Research. <https://doi.org/10.1162/153244303322533223>

Blei, David M., Ng, A. Y., & Jordan, M. I. (2003). Latent Dirichlet Allocation. Journal of Machine Learning Research. <https://doi.org/10.1162/jmlr.2003.3.4-5.993>

Cherenkova, Y., Zingaro, D., & Petersen, A. (2014). Identifying challenging CS1 concepts in a large problem dataset. In Proceedings of the 45th ACM technical symposium on Computer science education - SIGCSE '14. <https://doi.org/10.1145/2538862.2538966>

Cichocki, A., & Phan, A. H. (2009). Fast local algorithms for large scale nonnegative matrix and tensor factorizations. IEICE Transactions on Fundamentals of Electronics, Communications and Computer Sciences. <https://doi.org/10.1587/transfun.E92.A.708>

Corbett, A. T., & Anderson, J. R. (1995). Knowledge tracing: Modeling the acquisition of procedural knowledge. User Modeling and User-Adapted Interaction. <https://doi.org/10.1007/BF0109982>

Lee, D. D., & Seung, H. S. (1999). Learning the parts of objects by non-negative matrix factorization. Nature. <https://doi.org/10.1038/44565>

Mimno, D., Wallach, H. M., Talley, E., Leenders, M., & McCallum, A. (2011). Optimizing Semantic Coherence in Topic Models. In Proc. of the Conference on Empirical Methods in Natural Language Processing, EMNLP '11. <https://doi.org/10.1037/1082-989X.12.1.105>



References



Pavlik, P. I., Cen, H., & Koedinger, K. R. (2009). Performance factors analysis - A new alternative to knowledge tracing. In *Frontiers in Artificial Intelligence and Applications*.

<https://doi.org/10.3233/978-1-60750-028-5-531>

Petersen, A., Craig, M., & Zingaro, D. (2011). Reviewing CS1 exam question content. In *Proceedings of the 42nd ACM technical symposium on Computer science education - SIGCSE '11*.

<https://doi.org/10.1145/1953163.1953340>

Röder, M., Both, A., & Hinneburg, A. (2015). Exploring the Space of Topic Coherence Measures. In *Proceedings of the Eighth ACM International Conference on Web Search and Data Mining - WSDM '15*.

<https://doi.org/10.1145/2684822.2685324>

Sahebi, S., Lin, Y.-R., & Brusilovsky, P. (2016). Tensor Factorization for Student Modeling and Performance Prediction in Unstructured Domain. In *Educational Data Mining, proceedings of the 9th International Conference on*. http://www.educationaldatamining.org/EDM2016/proceedings/paper_150.pdf

Salton, G., & McGill, M. J. (1986). Introduction to modern information retrieval. *Introduction to Information Retrieval* (p. 400). McGraw-Hill, Inc. Retrieved from <http://portal.acm.org/citation.cfm?id=576628>

Sheard, J., Carbone, A., Clear, T., Raadt, M. De, Souza, D. D., Harland, J., ... Warburton, G. (2011). Exploring Programming Assessment Instruments : A Classification Scheme for Examination Questions. In *Seventh International Workshop on Computing Education Research*. <https://doi.org/10.1145/2016911.2016920>

References



Steyvers, M., & Griffiths, T. (2010). Probabilistic Topic Models. In Latent Semantic Analysis: A Road To Meaning. [https://doi.org/10.1016/s0364-0213\(01\)00040-4](https://doi.org/10.1016/s0364-0213(01)00040-4)

Zhang, W., Yoshida, T., & Tang, X. (2011). A comparative study of TF*IDF, LSI and multi-words for text classification. Expert Systems with Applications. <https://doi.org/10.1016/j.eswa.2010.08.066>

Yan, X., Guo, J., Liu, S., Cheng, X., & Wang, Y. (2012). Clustering short text using Ncut-weighted non-negative matrix factorization. In Proceedings of the 21st ACM international conference on Information and knowledge management - CIKM '12 (p. 2259). <https://doi.org/10.1145/2396761.2398615>